

Methods 3: Multilevel Statistical Modeling and Machine Learning

Week 06: Mid-way evaluation and Machine Learning Intro:
October 8, 2024

The course plan

Week 1: Introduction

Instructor sessions: *Setting up R and Python and recollection of the general linear model*

Week 2: Multilevel linear regression

Instructor sessions: *Modelling subject level effects – and how do they differ from group level effects?*

Week 3: Link functions and fitting generalised linear multilevel models

Instructor sessions: *What to do when the response variable is not continuous?*

Week 4: Evaluating Generalised linear mixed models

Instructor sessions: *How do we assess how models compare to one another?*

Week 5: Explanation and Prediction

Instructor sessions: *Code review*

Week 6: Mid-way evaluation and Machine Learning Intro

Instructor sessions: *Getting Python Running*

Week 7: Linear regression revisited (machine learning)

Instructor sessions: *How to constrain our models to make them more predictive*

Week 8: Logistic regression revisited (machine learning)

Instructor sessions: *Categorizing responses based on informed guesses*

Week 9: Dimensionality Reduction, Principled Component Analysis (PCA)

Instructor sessions: *What to do with very rich data?*

Week 10: Outlook, unsupervised classification and neural networks

Instructor sessions: *Data with no labels and networks*

Week 11: Organising and preprocessing messy data

Instructor sessions: *Code review*

Week 12: Final evaluation and wrap-up of course

Instructor sessions: *Ask anything!*

The four classical levels of variables

- Nominal
 - examples: true/false, correct/incorrect, female/male, dog/cat, apples/pear, also called *categorical*
 - they are **names** of categories, but it does not make sense to order them
- Ordinal
 - examples: senior/junior, adult/child 1/2/3
 - they are also **names** of categories, but there is an explicit or implicit **ordering**, i.e. one is greater than another
- Interval
 - examples: the year 1984 AD; the temperature 100 °C
 - they can be **continuous**, and there is **ordering**, e.g. 100 °C > 90 °C. And intervals can be compared, e.g. the interval from 80 °C to 100 °C is as long as the one from 40 °C to 60 °C
 - crucially, there is no real 0; the year before 1 AD is not characterised by absence of time; and 0 °C is not characterised by the absence of temperature. This means that we cannot say that, say, 40 °C is twice as high a temperature as 20 °C
- Ratio
 - examples: the temperature 273 K, the reaction time of a subject
 - they can be **continuous**, there is **ordering** and there is a **real 0**.
 - Thus we can say that 200 K is twice the temperature of 100 K, as 0 K *is* the absence of temperature; and we can say that subject 2, 400 ms, is twice as fast as subject 1, 200 ms, because 0 ms *is* the time when the event happened

Overview – pooling

- Complete pooling
 - Ignores the categorical predictor, e.g. *Subject*, *altogether*
 - $\text{lm}(\text{Reaction} \sim \text{Days})$
- No pooling
 - Overfits the categorical predictor, e.g. *Subject*, i.e. overstates the variation among *Subjects*
 - $\text{lm}(\text{Reaction} \sim \text{Days} * \text{Subject} - 1)$; *models slopes and intercepts for each subject*
 - $\text{lm}(\text{Reaction} \sim \text{Days} + \text{Subject} - 1)$; *models intercepts for each subject*
- Partial pooling
 - A compromise between the two extremes above. If a group, e.g. *Subject*, has few observations (high variance), it will be shrunk towards the overall mean. If *Subject* has many observations (low variance), it will be shrunk less towards the overall mean
 - $\text{lmer}(\text{Reaction} \sim \text{Days} + (\text{Days} | \text{Subject})$ # *models slopes and intercepts for each subject*
 - $\text{lmer}(\text{Reaction} \sim \text{Days} + (1 | \text{Subject})$ # *models intercepts for each subject*

Recap?

- Bias can be added in ways that improve prediction
 - it does this by removing collinearity
 - thereby making the model stable
 - and more generalisable

Learning goals and outline

Mid-way evaluation and Machine Learning Intro

- 1) Learning some early *classification* methods
 - Perceptron and ADALine
 - Classification depends on having a quantiser function
- 2) Learning how linear *regression* (with biasing penalties) can be constructed and cross-validated

Mid-way evaluation

~10 min


- 1) Write something you liked about the course so far
- 2) Write something you did not like about the course so far
- 3) What would you change?

I'll summarise the feedback on the three points, and what we'll change for next time

The Perceptron

Raschka S (2015) Python Machine Learning. Packt Publishing Ltd

Black box idea



$$\mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$$

$$z = w_1 x_1 + \dots + w_m x_m$$

Question: what do x , w and z correspond to in the above picture of the *Perceptron*?

(p. 18: Raschka, 2015)

Prediction/classification rule

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ -1 & \text{otherwise} \end{cases}$$

Perceptron fires

Perceptron doesn't fire

θ is a pre-specified threshold

Prediction/classification rule

$$w_0 = -\theta$$

$$x_0 = 1$$

$$z = w_0 x_0 + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$

$$z = -\theta + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$


$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

Perceptron fires

Perceptron doesn't fire

Perceptron classification

We want to find $\mathbf{w}^T \mathbf{x}$ that achieves this separation



In *Python* (in 2021)

```
class Perceptron(object):  
    """ Perceptron classifier  
  
    Parameters  
    -----  
    eta : float  
        Learning rate (between 0.0 and 1.0)  
    n_iter : int  
        Passes over the training dataset.  
  
    Attributes  
    -----  
    w_ : 1d-array  
        Weights after fitting.  
    errors_ : list  
        Number of misclassifications in every epoch.  
  
    """
```

Special definition that indicates what the object (*Perceptron*) can be initialised with

```
def __init__(self, eta=0.01, n_iter=10):  
    self.eta = eta  
    self.n_iter = n_iter
```

```
ppn = Perceptron(eta=0.1, n_iter=10)
```

Specifying methods of *Perceptron*


1. Initialize the weights to 0 or small random numbers.
2. For each training sample $\mathbf{x}^{(i)}$ perform the following steps:
 1. Compute the output value \hat{y} .
 2. Update the weights.

```
def fit(self, X, y):  
    """ Fit training data.  
  
    Parameters  
    -----  
    X : {array-like}, shape = [n_samples, n_features]  
        Training vectors, where n_samples  
        is the number of samples and  
        n_features is the number of features.  
    y : array-like, shape = [n_samples]  
        Target values.  
  
    Returns  
    -----  
    self : object  
  
    """  
    self.w_ = np.zeros(1 + X.shape[1])  
    self.errors_ = []  
  
    for _ in range(self.n_iter):  
        errors = 0  
        for xi, target in zip(X, y):  
            update = self.eta * (target - self.predict(xi))  
            self.w_[1:] += update * xi  
            self.w_[0] += update  
            errors += int(update != 0.0)  
        self.errors_.append(errors)  
    return self
```


Compute the output value \hat{y}



```
def net_input(self, X):  
    """Calculate net input"""  
    return np.dot(X, self.w_[1:]) + self.w_[0]  
  
def predict(self, X):  
    """Return class label after unit step"""  
    return np.where(self.net_input(X) >= 0.0, 1, -1)
```



$$\phi(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ -1 & \text{otherwise} \end{cases}$$

When we are right

$$\Delta w_j = \eta (\overset{\text{real label}}{\boxed{-1}} - \underset{\text{predicted label}}{\boxed{-1}}) x_j^{(i)} = 0$$

$$\Delta w_j = \eta (\overset{\text{real label}}{\boxed{1}} - \underset{\text{predicted label}}{\boxed{1}}) x_j^{(i)} = 0$$

```
update = self.eta * (target - self.predict(xi))  
self.w_[1:] += update * xi  
self.w_[0] += update
```

Δw_j : change in weight

η : learning rate

When we are wrong


$$\Delta w_j = \eta (\overset{\text{real label}}{\boxed{1}} - \underset{\text{predicted label}}{\boxed{-1}}) x_j^{(i)} = \eta (2) x_j^{(i)}$$

$$\Delta w_j = \eta (\overset{\text{real label}}{\boxed{-1}} - \underset{\text{predicted label}}{\boxed{1}}) x_j^{(i)} = \eta (-2) x_j^{(i)}$$

Δw_j : change in weight

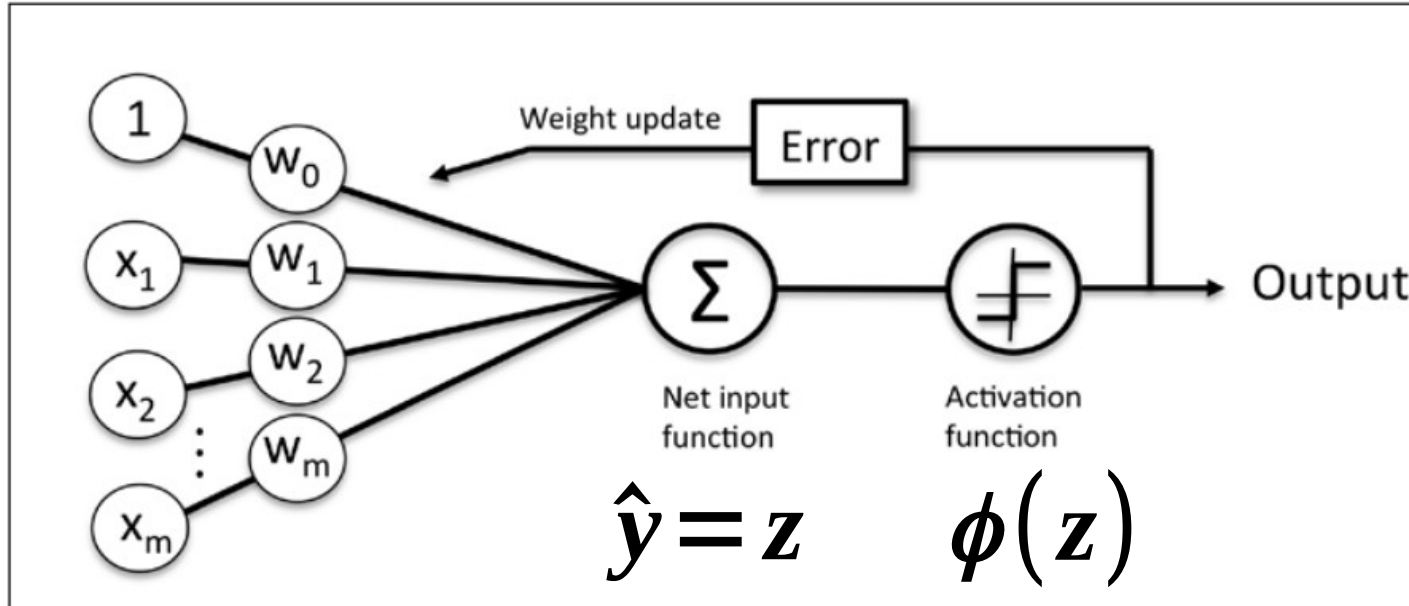
η : learning rate

Convergence only possible when linearly separable




(p. 23: Raschka, 2015)

Perceptron: Graphical summary




(p. 24: Raschka, 2015)

An example



```
In [153]: ppn.w_  
Out[153]: array([-0.04 , -0.068,  0.182])
```




```
In [155]: X_subset[0, :]  
Out[155]: array([5.1, 1.4])
```

$$\hat{y}_1 = -0.04 - 0.068 \cdot 5.1 + 0.182 \cdot 1.4 = -0.132$$

```
In [156]: ppn.net_input(X_subset[0, :])  
Out[156]: -0.13199999999999953
```

(p. 29 & p. 32: Raschka, 2015)

Perceptron: Graphical summary




(p. 24: Raschka, 2015)

ADALINE

Raschka S (2015) Python Machine Learning. Packt Publishing Ltd

ADaptive Linear NEuron (ADALINE)



$$w^T x = w_0 x_0 + w_1 x_1 + \dots + w_{m-1} x_{m-1} + w_m x_m$$

ADALINE Gradient descent

```
def __init__(self, eta=0.01, n_iter=50):  
    self.eta = eta  
    self.n_iter = n_iter
```

```
class AdalineGD(object):  
    """ ADaptive LInear NEuron classifier  
  
    Parameters  
    -----  
    eta : float  
        Learning rate (between 0.0 and 1.0)  
    n_iter : int  
        Passes over the training dataset.  
  
    Attributes  
    -----  
    w_ : 1d-array  
        Weights after fitting.  
    errors_ : list  
        Number of misclassifications in every epoch.  
  
    """
```

Methods

$$\hat{y} = w^T x$$

$$z(w^T x) = w^T x$$

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

```
def net_input(self, X):  
    """Calculate net input"""  
    return np.dot(X, self.w_[1:]) + self.w_[0]  
  
def activation(self, X):  
    """Computer linear activation"""  
    return self.net_input(X)  
  
def predict(self, X):  
    """Return class label after unit step"""  
    return np.where(self.activation(X) >= 0.0, 1, -1)
```

The *fit* method

eta (η): learning rate (a constant)

output: $X\hat{\beta} = \hat{y}$

errors: $y - \hat{y}$

$X.T.dot(errors)$: $X^T \cdot (y - \hat{y})$

$X^T \cdot (y - \hat{y}) = \Delta w_1 + \Delta w_2 + \dots + \Delta w_{m-1} + \Delta w_m$

cost function: $(\sum (y - \hat{y})^2) / 2$

```
def fit(self, X, y):
    """ Fit training data.

    Parameters
    -----
    X : {array-like}, shape = [n_samples, n_features]
        Training vectors, where n_samples
        is the number of samples and
        n_features is the number of features.
    y : array-like, shape = [n_samples]
        Target values.


    Returns
    -----
    self : object

    """
    self.w_ = np.zeros(1 + X.shape[1])
    self.cost_ = []

    for i in range(self.n_iter):
        output = self.net_input(X)
        errors = (y - output)
        self.w_[1:] += self.eta * X.T.dot(errors)
        self.w_[0] += self.eta * errors.sum()
        cost = (errors**2).sum() / 2.0
        self.cost_.append(cost)
    return self
```

Learning rate

Overshooting
the global
minimum




Slow stepping
down

Always check whether it converges

(p. 40: Raschka, 2015)


Gradient descent



$$J(w) = (\sum (y - \hat{y})^2) / 2$$

(p. 40: Raschka, 2015)

Adaline - Gradient Descent



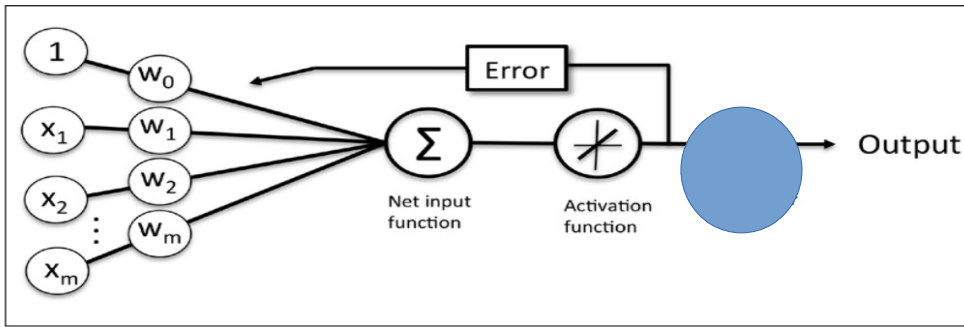
(p. 42: Raschka, 2015)

```
In [163]: ada.w_  
Out[163]: array([-0.40808285, -0.33924452,  0.79202224])
```

```
In [155]: X_subset[0, :]  
Out[155]: array([5.1, 1.4])
```

$$\hat{y}_1 \approx -0.408 - 0.339 \cdot 5.1 + 0.792 \cdot 1.4 = -1.03$$

```
In [164]: ada.net_input(X_subset[0, :])  
Out[164]: -1.029398778794785
```

How does this relate to linear regression?


From ADALINE

$$\mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + \dots + w_{m-1} x_{m-1} + w_m x_m$$

Very similar to ADALINE
and when converged will be virtually
identical to the ordinary least
squares solution

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

Convergence



```
class LinearRegressionGD(object):  
  
    def __init__(self, eta=0.001, n_iter=20):  
        self.eta = eta  
        self.n_iter = n_iter  
  
    def fit(self, X, y):  
        self.w_ = np.zeros(1 + X.shape[1])  
        self.cost_ = []  
  
        for i in range(self.n_iter):  
            output = self.net_input(X)  
            errors = (y - output)  
            self.w_[1:] += self.eta * X.T.dot(errors)  
            self.w_[0] += self.eta * errors.sum()  
            cost = (errors**2).sum() / 2.0  
            self.cost_.append(cost)  
        return self  
  
    def net_input(self, X):  
        return np.dot(X, self.w_[1:]) + self.w_[0]  
  
    def predict(self, X):  
        return self.net_input(X)
```

ADALINE

```
def fit(self, X, y):
    """ Fit training data.

    Parameters
    -----
    X : {array-like}, shape = [n_samples, n_features]
        Traing vectors, where n_samples
        is the number of samples and
        n_features is the number of features.
    y : array-like, shape = [n_samples]
        Target values.

    Returns
    -----
    self : object

    """
    self.w_ = np.zeros(1 + X.shape[1])
    self.cost_ = []

    for i in range(self.n_iter):
        output = self.net_input(X)
        errors = (y - output)
        self.w_[1:] += self.eta * X.T.dot(errors)
        self.w_[0] += self.eta * errors.sum()
        cost = (errors**2).sum() / 2.0
        self.cost_.append(cost)
    return self
```

```
def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]

def activation(self, X):
    """Computer linear activation"""
    return self.net_input(X)

def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.activation(X) >= 0.0, 1, -1)
```

```
class LinearRegressionGD(object):
```

```
    def __init__(self, eta=0.001, n_iter=20):
        self.eta = eta
        self.n_iter = n_iter

    def fit(self, X, y):
        self.w_ = np.zeros(1 + X.shape[1])
        self.cost_ = []

        for i in range(self.n_iter):
            output = self.net_input(X)
            errors = (y - output)
            self.w_[1:] += self.eta * X.T.dot(errors)
            self.w_[0] += self.eta * errors.sum()
            cost = (errors**2).sum() / 2.0
            self.cost_.append(cost)
        return self

    def net_input(self, X):
        return np.dot(X, self.w_[1:]) + self.w_[0]

    def predict(self, X):
        return self.net_input(X)
```

LINEAR REGRESSION



Olivier Grisel
@ogrisel

In scikit-learn 1.0, we decided to deprecate the `sklearn.datasets.load_boston` function because the design of this dataset casually assumes that people prefer to buy housing in racially segregated neighborhoods.

(p. 229: Raschka, 2015)

The features of the 506 samples may be summarized as shown in the excerpt of the dataset description:

- **CRIM**: This is the per capita crime rate by town
- **ZN**: This is the proportion of residential land zoned for lots larger than 25,000 sq.ft.
- **INDUS**: This is the proportion of non-retail business acres per town
- **CHAS**: This is the Charles River dummy variable (this is equal to 1 if tract bounds river; 0 otherwise)
- **NOX**: This is the nitric oxides concentration (parts per 10 million)
- **RM**: This is the average number of rooms per dwelling
- **AGE**: This is the proportion of owner-occupied units built prior to 1940
- **DIS**: This is the weighted distances to five Boston employment centers
- **RAD**: This is the index of accessibility to radial highways
- **TAX**: This is the full-value property-tax rate per \$10,000
- **PTRATIO**: This is the pupil-teacher ratio by town
- **B**: This is calculated as $1000(B_k - 0.63)^2$, where B_k is the proportion of people of African American descent by town
- **LSTAT**: This is the percentage lower status of the population
- **MEDV**: This is the median value of owner-occupied homes in \$1000s

California Housing Dataset

x

MedInc: median income in block group

HouseAge: median house age in block group

AveRooms: average number of rooms per household

AveBedrms: average number of bedrooms per household

Population: block group population


AveOccup: average number of household members

Latitude: block group latitude

Longitude: block group longitude

y

MedHouseVal: The median house value for California districts, expressed in hundreds of thousands of dollars (\$100,000)



```
LR = LinearRegressionGD()
LR.fit(X[:, 0:1], y) ## just fitting on Median Income
```

```
print(LR.w_)
```

```
## [-1.10873027e+51 -5.27207344e+51]
```

!


CC BY Licence 4.0: Lau Møller Andersen

$$J(w) = (\sum (y - \hat{y})^2) / 2$$

Check for convergence!

```
X, y = fetch_california_housing(return_X_y=True)
from sklearn.preprocessing import StandardScaler
sc_x = StandardScaler()
sc_y = StandardScaler()
X_std = sc_x.fit_transform(X)
y_std = np.squeeze(sc_y.fit_transform(y.reshape(-1, 1)))

LR = LinearRegressionGD(eta=1e-6, n_iter=1000)
LR.fit(X_std[:, 0:1], y_std) ## just fitting on Median Income
```



```
class LinearRegressionGD(object):
    def __init__(self, eta=0.001, n_iter=20):
        self.eta = eta
        self.n_iter = n_iter
```


- Standardisation

$$x_{std}^{(i)} = \frac{x^{(i)} - \mu_x}{\sigma_x}$$

μ_x : sample mean for the feature: x


σ_x : sample standard deviation for the feature: x

Brings data onto a normal distribution with $\mu=0$ and $\sigma=1$



Multiple linear regression

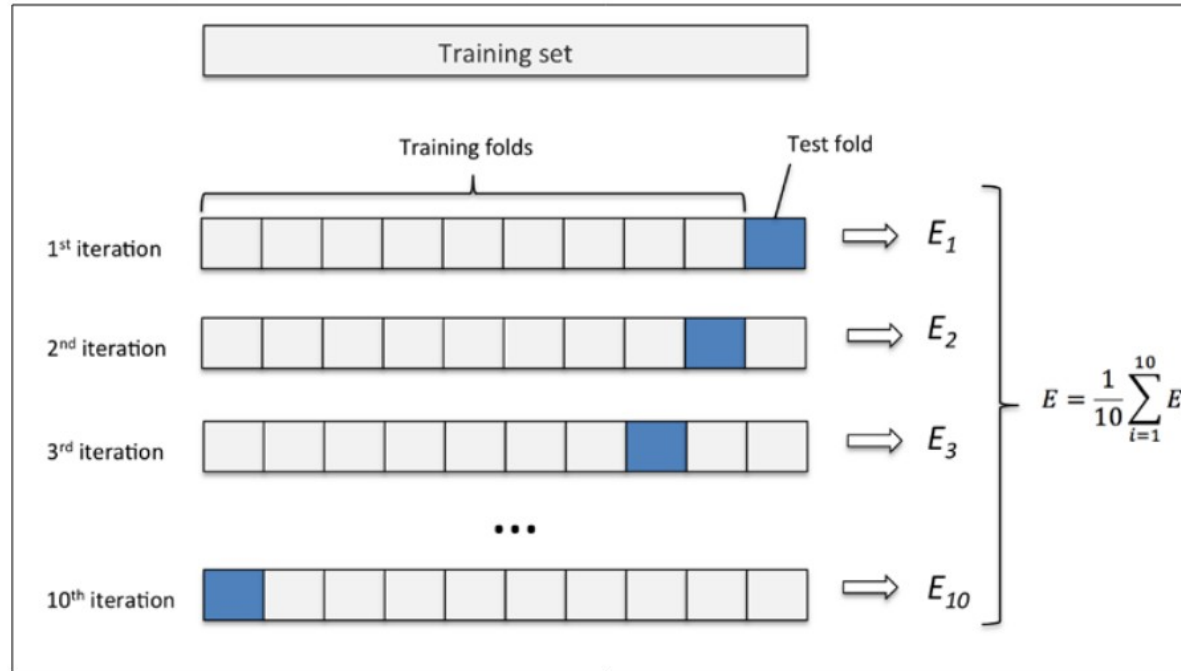
$$\begin{aligned} \hat{MedHouseVal} = & \hat{w}_0 x_0 + \hat{w}_1 MedInc + \hat{w}_2 HouseAge \\ & + \hat{w}_3 AveRooms + \hat{w}_4 AveBedrms \\ & + \hat{w}_5 Population + \hat{w}_6 AveOccup + \hat{w}_7 Latitude + \hat{w}_8 Longitude + \epsilon \end{aligned}$$



Because of the collinearity, we know we are prone to overfitting, so we do **out-of-sample** prediction instead of validating our model with traditional measures like R^2 and maximum likelihood

How to choose the **out-of-sample** dataset?


Cross-validation



(p. 176: Raschka, 2015)

CC BY Licence 4.0: Lau Møller Andersen

```
OLS = LinearRegression()  
OLS.fit(X_std, y_std)  
  
MSE = np.mean(cross_validate(OLS, X_std, y_std, k=11))
```




We can impose penalties


(but not on the intercept)

$$J(w)_{\text{Ridge}} = \sum^n \left(y^{(i)} - \hat{y}^{(i)} \right)^2 + \lambda \|w\|_2^2$$

Ridge Regression: min MSE (0.3607) at: $\lambda = 51.7$





Ridge: MSE = 0.3607: 11-fold cross-validation. Variance = 0.6532



We can impose penalties

(but not on the intercept)

$$J(w)_{LASSO} = \sum_{i=1}^n \left(y^{(i)} - \hat{y}^{(i)} \right)^2 + \lambda \|w\|_1$$



Coefficients are shrunk

```
In [131]: OLS.coef_
```

```
Out[131]:
```

```
array([-0.09874812,  0.12473758,  0.02386168,  0.06945318, -0.22231612,  
        0.2911837 ,  0.0325356 , -0.32907266,  0.34212986, -0.28361575,  
       -0.21482076,  0.09763631, -0.43131412])
```

```
In [132]: RR.coef_
```

```
Out[132]:
```

```
array([-0.07536542,  0.08180161, -0.03182673,  0.07855868, -0.13185121,  
        0.31082552,  0.00548938, -0.23491485,  0.11242408, -0.0959682 ,  
       -0.18436614,  0.09154233, -0.36579723])
```

```
In [133]: lasso.coef_
```

```
Out[133]:
```

```
array([-0.07348948,  0.09107936, -0.          ,  0.07003181, -0.17110318,  
        0.30950805,  0.          , -0.29010247,  0.16456993, -0.1319311 ,  
       -0.19668957,  0.09063304, -0.41664587])
```

Summary

- We can build simple classification tools using scikit-learn
 - These can give us decision boundaries
 - That we can apply to new data (we haven't done that yet)
- We need to define cost functions
 - These can be conceptually separated from prediction functions (remember ADALINE)
- We can use linear regression to do continuous predictions

The course plan

Week 1: Introduction

Instructor sessions: *Setting up R and Python and recollection of the general linear model*

Week 2: Multilevel linear regression

Instructor sessions: *Modelling subject level effects – and how do they differ from group level effects?*

Week 3: Link functions and fitting generalised linear multilevel models

Instructor sessions: *What to do when the response variable is not continuous?*

Week 4: Evaluating Generalised linear mixed models

Instructor sessions: *How do we assess how models compare to one another?*

Week 5: Explanation and Prediction

Instructor sessions: *Code review*

Week 6: Mid-way evaluation and Machine Learning Intro

Instructor sessions: *Getting Python Running*

Week 7: Linear regression revisited (machine learning)

Instructor sessions: *How to constrain our models to make them more predictive*

Week 8: Logistic regression revisited (machine learning)

Instructor sessions: *Categorizing responses based on informed guesses*

Week 9: Dimensionality Reduction, Principled Component Analysis (PCA)

Instructor sessions: *What to do with very rich data?*

Week 10: Outlook, unsupervised classification and neural networks

Instructor sessions: *Data with no labels and networks*

Week 11: Organising and preprocessing messy data

Instructor sessions: *Code review*

Week 12: Final evaluation and wrap-up of course

Instructor sessions: *Ask anything!*

Learning goals and outline

Mid-way evaluation and Machine Learning Intro

- 1) Learning some early *classification* methods
 - Perceptron and ADALine
 - Classification depends on having a quantiser function
- 2) Learning how linear *regression* (with biasing penalties) can be constructed and cross-validated

Next time

- We will do logistic regression and linear regression together
 - but we will skip an assignment and put in a code review instead
 - P7
 - ~~Assignment 3: How to constrain our models to make them more predictive~~
 - **Assignment 4: Using logistic regression to classify subjective experience from brain data**
 - P8
 - code review
- I'll update the syllabus accordingly

Reading questions

- Chapter 10
 - What is lasso and ridge regression?
 - What is standardisation?
- Chapter 3
 - What is the cost function of logistic regression?
 - What is overfitting and underfitting?
 - What is a support vector?